


Oracle Database JSON Workshop Preview

Josh Spiegel

Software Architect, Oracle Database



Oracle Database as a JSON Document Store



- A type of NoSQL database
- Low-latency, high throughput
- Store collections of JSON objects
 - Simple
 - Flexible
 - App friendly
- Low-cost

```
{  
  "movie_id" : 1652,  
  "title" : "Iron Man 2",  
  "date" : "2010-05-07",  
  "cast" : [  
    "Robert Downey Jr.",  
    "Larry Ellison"  
  ]  
}
```

Today's JSON Workshop

Introduction


- Benefits of JSON document databases
- Oracle Database as a document store

SQL/JSON

- JSON Storage
- Query and update JSON
- Load JSON from files

JSON Collections

- Underlying data model
- Connecting Mongo DB clients to Oracle Database
- Migrating data from MongoDB to Oracle



First demo of new and unreleased features!

Demonstration


Oracle Cloud Infrastructure – **Ashburn Data Center**



Autonomous JSON
Database 19c



Oracle Database
23c Free



Oracle REST Data
Services **23.2** RC1

Linux Compute Node

Oracle Database JSON


Josh Spiegel

Software Architect, Oracle Database




Oracle Converged Database

Converged Database Architecture



for **any** data type or workload

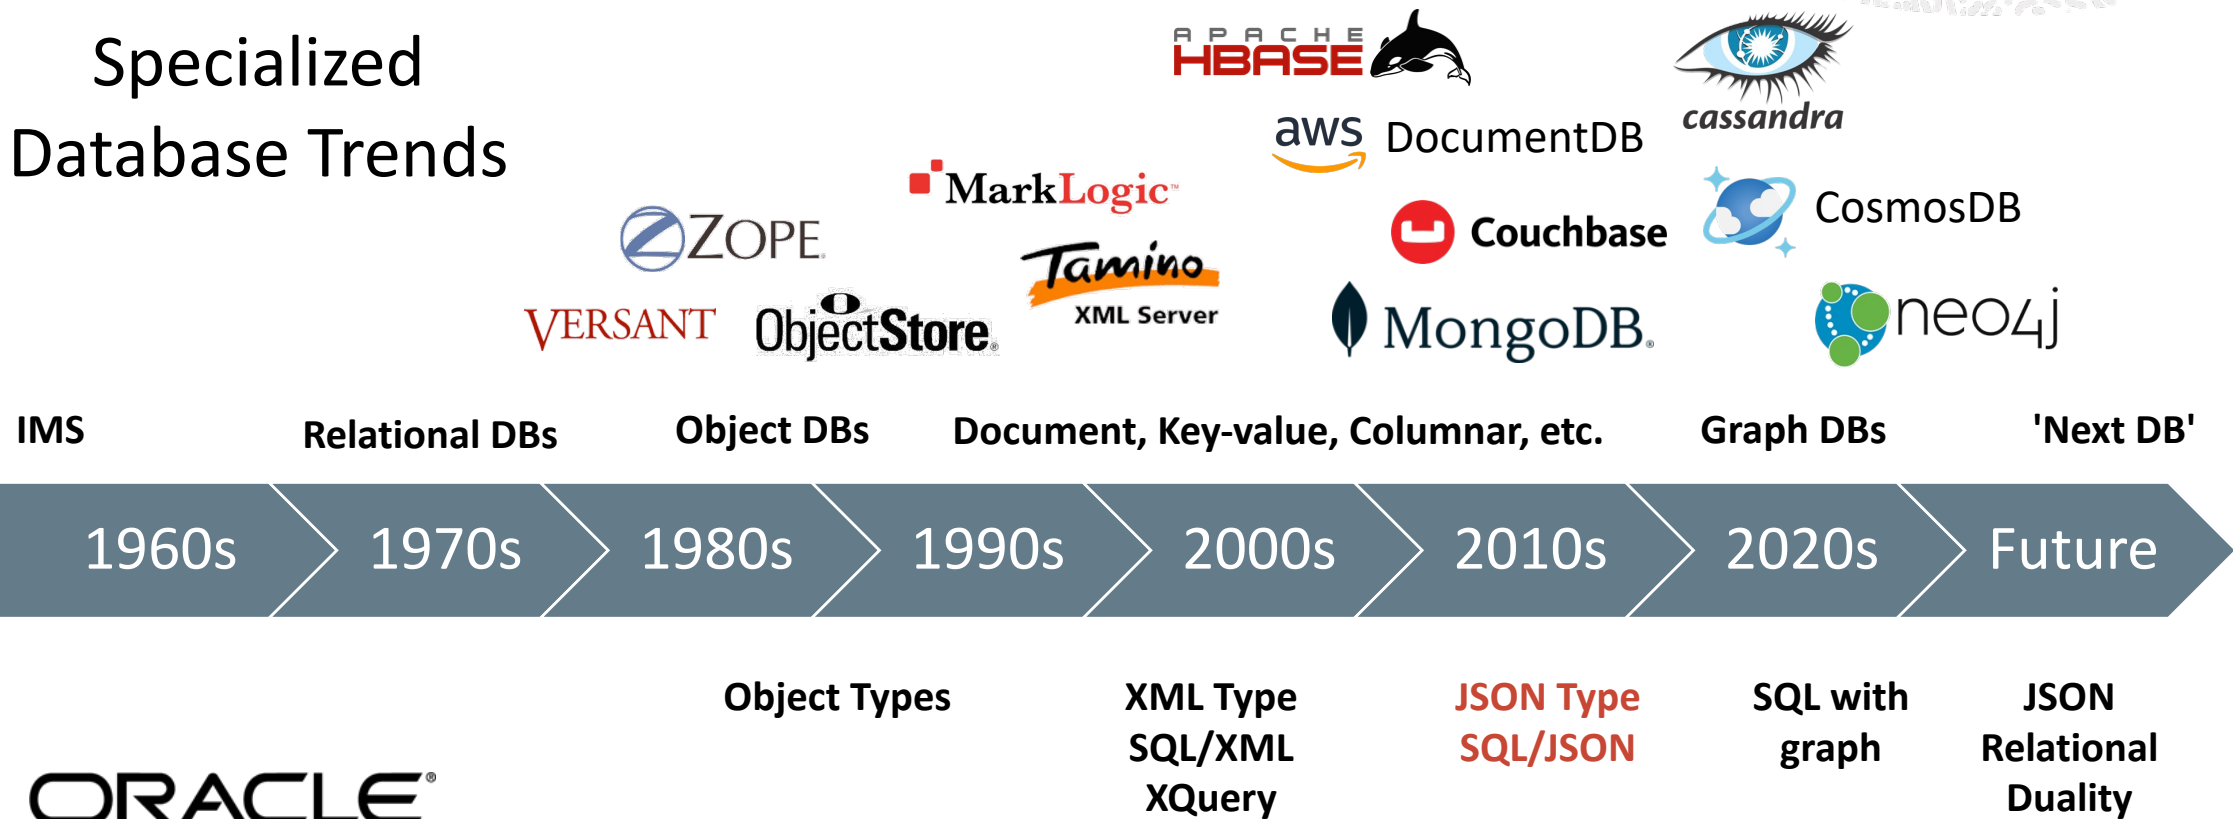
Single-purpose databases



for each data type and workload
Multiple security models, languages, skills, licenses, etc

Specialized vs Multi-Model

Specialized Database Trends



ORACLE®
Multi-Model SQL




Database Popularity 2023 (db-engines.com)

Top-10 Ranked Databases


Database	Type	Popularity
Oracle	Relational, Multi-model	1231.48
MySQL	Relational, Multi-model	1163.94
Microsoft SQL Server	Relational, Multi-model	930.06
PostgreSQL	Relational, Multi-model	612.82
MongoDB	Document	425.36
Redis	Key-value, Multi-model	167.35
IBM Db2	Relational, Multi-model	144.89
Elasticsearch	Search engine, Multi-model	143.75
Microsoft Access	Relational	134.45
SQLite	Relational	131.21

Total Popularity % by Type



© 2023, DB-Engines.com

NoSQL JSON Document Stores




- Store collections of JSON (JavaScript Object Notation)
- A simple, self-describing, flexible way to model objects
- String, numbers, true/false, Nested repeating values

```
{  
  "movie_id" : 1652,  
  "title" : "Iron Man 2",  
  "date" : "2010-05-07",  
  "cast" : [  
    "Robert Downey Jr.",  
    "Larry Ellison"  
  ]  
}
```



Why store JSON in a database?



```
{
  "movie_id" : 1652,
  "title" : "Iron Man 2",
  "date" : "2010-05-07",
  "cast" : [
    "Robert Downey Jr.",
    "Larry Ellison",
    ..
  ]
}
```

Schema-flexible

- No upfront schema design
- Application-controlled schema
- Simple data model

```
class Movie {

    int movie_id;
    String title;
    LocalDateTime date;
    List<String> cast;

    Movie() {
        ...
    }
}
```


Less Impedance Mismatch

- Maps to application objects
- Supports nested structures
- Read/write without joins



Data Exchange

- One representation across tiers
- Human readable
- REST, JSON are ubiquitous



NoSQL Document Store Features



Elastic compute
and storage



Single-digit latency
reads and writes



Highly available



Low-cost

NoSQL Document Store - Data as code

NoSQL Access


```
db.movies.insertOne(movieValue);  
  
db.movies.find({"_id" : 123});
```

SQL Access

```
PreparedStatement ins =  
    con.prepareStatement( "INSERT INTO movies VALUES (:1,:2:)" );  
  
ins.setNumber(1, id)  
ins.setString(2, movieName);  
ins.execute();  
  
ResultSet rs =  
    stmt.executeQuery( "SELECT * FROM movies WHERE id = 123" )
```



Oracle Database: all the features of a single-purpose JSON document store – with superior performance



** Autonomous JSON Database with 8 OCPU running in San Jose region, Mongo API





But without losing all the strengths of Oracle Database when you store JSON...

- ANSI SQL / JSON
- Advanced analytics
- In-memory columnar
- Full-text search
- ACID transactions
- Consistency and durability
- Mission critical use cases
- Secure by default
- Resource management

Autonomous Database offers workload-optimized options



OLTP+ JSON workloads

Autonomous Database for Transaction Processing and mixed workloads - ATP

Pre-configured for row format, indexes, and data caching to accelerate **transaction processing** and **mixed workloads**

- **Converged database** with no data storage limits (relational or JSON), JSON data fully supported
- Includes **Oracle Database API for MongoDB**

Single-click upgrade to ATP



JSON-centric workloads

- Autonomous JSON Database – AJD

Price optimized for transactions and analytics on **JSON data**

- Same features as ATP but **75% lower price**
- **Unlimited JSON collections** + up to 20GB non-JSON data
- Includes **Oracle Database API for MongoDB**
- Great option for a “start small and grow big” app dev strategy—can be **upgraded to ATP with a click of a button**




Analytics workloads including JSON data

- Autonomous Database for analytics and data warehousing – ADW

Pre-configured for columnar format, partitioning, and large joins to accelerate **analytics, data warehouse, and data lakehouse**

- Same features as ATP, optimized for analytics
- JSON data fully supported
- Includes **Oracle Database API for MongoDB**



Autonomous JSON Database



Elastic compute
and storage



Single-digit latency
reads and writes



Highly available



Low-price,
always-free tier

Accessing JSON with SQL



Traditional Relational Model

- A schema contains tables
- A table contains rows
- A table is flat
- Rows are structured
- Data is accessed with SQL
- Related rows are joined

Flat, structured tables

id	title	date
123	Iron Man	2010-05-07
345	Thor	2022-07-08

Accessed with SQL

```
SELECT m.title, m.date
FROM   movies m
WHERE  m.id = 123
```

Relational Model + JSON Columns

- Schema-flexible JSON stored within a structured column
- SQL extended to process JSON column values
- Stored using query-efficient OSON binary format

```
CREATE TABLE movies (data JSON);

INSERT INTO movies VALUES (
  JSON {
    'id'      : 123,
    'title'   : 'Iron Man'
  }
);
```

JSON Storage History

12c

JSON-text storage
and query processing
(clob, blob, varchar2)

19c

Binary JSON storage (OSON)
and Mongo API support added
for **Autonomous** Databases
(in BLOB columns)

18c

21c

Native JSON datatype and
collections backed by OSON
(all database types)

Why OSON? - Extended types

- Standard
 - OBJECT { }
 - ARRAY []
 - STRING
 - TRUE/FALSE
 - NULL
 - NUMBER
- Extended
 - BINARY_FLOAT
 - BINARY_DOUBLE
 - TIMESTAMP/DATE
 - INTERVALS/INTERVALYM
 - RAW

Fidelity with relational data


```
CREATE TABLE orders VALUES (  
    oid          NUMBER,  
    created      TIMESTAMP,  
    status       VARCHAR2(10),  
);  
  
{  
  "oid":123,  
  "created":"2020-06-04T12:24:29Z",  
  "status":"OPEN"  
}
```



Why OSON? - Performance

- Faster path evaluation
- Efficient random access
- Smaller than JSON text and BSON
- Less network and storage IO
- No text parsing or serialization

Yahoo Cloud Service Benchmark



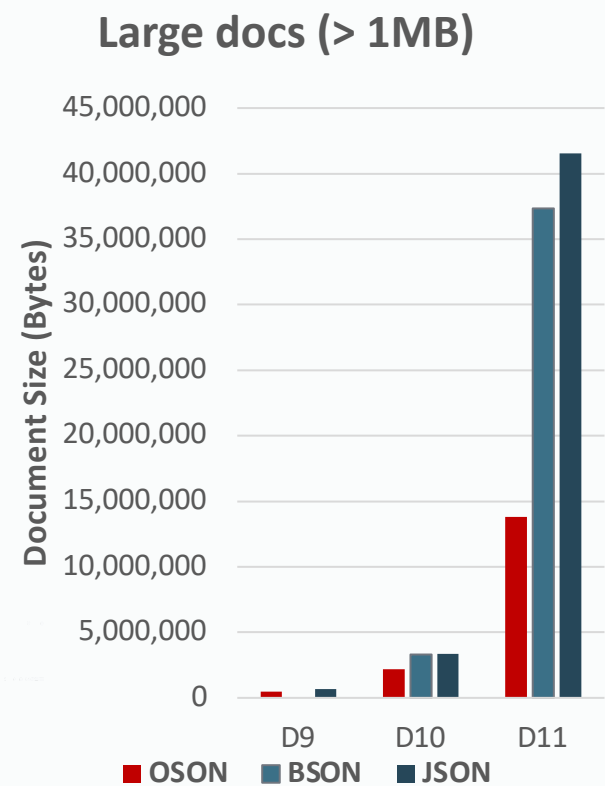


See VLDB 2020:

Native JSON Datatype Support: Maturing SQL and NoSQL convergence in Oracle Database



OSON Compression

Up to 2.7x
more compact

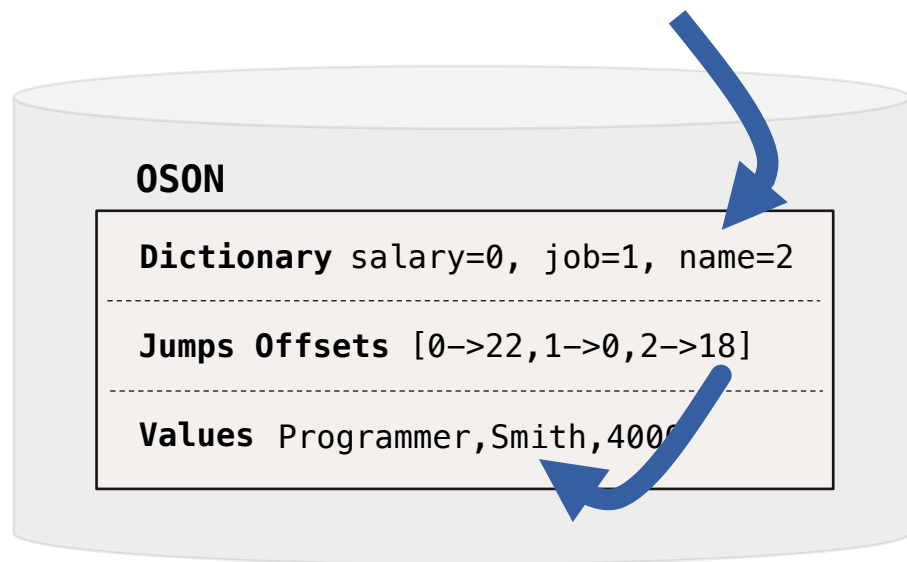


OSON Performance


```
{"name": "Smith", "job": "Programmer", "salary": 40000}
```

SQL

```
e.data.name.string()
```



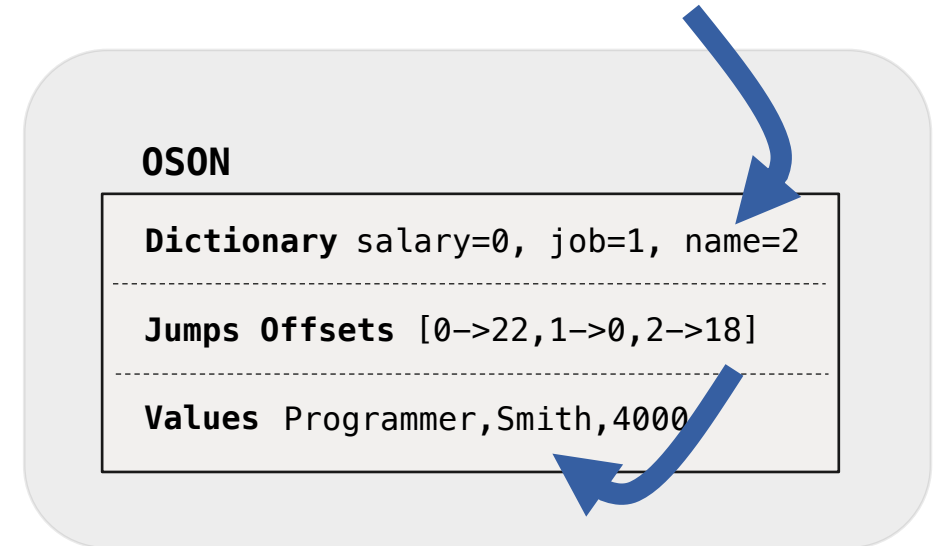
Database



Network

Java

```
obj.getString("name")
```



Application



SQL/JSON

- Use SQL to query JSON data
 - JSON to relational
 - Relational to JSON
- Joins, aggregation, projection
- Construct new JSON values
- Update JSON values
- Unnest arrays

JSON aggregation and construction

```
SELECT JSON {  
  'total' : sum(t.jcol.gross.number()),  
  'genre' : m.jcol.genre.string()  
}  
FROM movies m  
GROUP BY e.jcol.genre.string()
```

JSON unnesting

```
SELECT title, actor  
FROM movies NESTED jcol COLUMNS (  
  title,  
  NESTED actors[*] COLUMNS (  
    actor  
  )  
)  
WHERE id = 123
```

Simple Dot Notation

```
CREATE TABLE movies(  
    id        number PRIMARY KEY,  
    title     varchar2(50),  
    year      number,  
    data      JSON);  
  
INSERT INTO movies VALUES( 1, 'Avatar', 2009,  
    '{"sku": "LYG56160", "runtime": 162,  
    "cast": ["Sam Worthington", "Zoe Saldana"],  
    "studio": ["20th Century Studios",  
              "Lightstorm Entertainment"] }'  
);
```

```
SELECT m.id,  
       m.title,  
       m.data.sku,  
       m.data.genre,  
       m.data.studio  
FROM movies m  
WHERE m.data.runtime > 120;
```

Select JSON attributes as if they were columns.



SQL | Oracle Database Actions

ppkhnzjhg74axsq-ajd21cdb.adb.eu-frankfurt-1.oraclecloudapps.com/ords/admin/_sdw/?nav=worksheet

Search

ADMIN

Navigator

Files

ADMIN

Tables

Search...

MOVIES

- ID
- TITLE
- YEAR
- DATA

[Worksheet]*

Consumer group: LOW

Data Load

```
1 SELECT id,
2       title,
3       m.data.sku,
4       m.data.genre,
5       m.data.studio,
6       m.data.runtime
7 FROM movies m
8 WHERE m.data.runtime.number() > 120;
9
```

Query Result

Script Output

DBMS Output

Explain Plan

Autotrace

SQL History

Download

4:00:52 AM - Code execution finished.

Powered by ORDS



JSON_QUERY

```
CREATE TABLE movies(  
    id        number PRIMARY KEY,  
    title     varchar2(50),  
    year      number,  
    data      JSON);  
  
INSERT INTO movies VALUES( 1, 'Avatar', 2009,  
    '{"sku": "LYG56160", "runtime": 162,  
    "cast": ["Sam Worthington", "Zoe Saldana"],  
    "studio": ["20th Century Studios",  
              "Lightstorm Entertainment"] }'  
);
```

```
SELECT  
    id,  
    title,  
    m.data.sku,  
    JSON_QUERY(  
        m.data,  
        '$.genre'  
        EMPTY ARRAY ON EMPTY  
    ),  
    m.data.studio  
FROM movies m  
WHERE m.data.runtime > 120;
```

JSON_QUERY to have more control over how a JSON value read.

JSON_TABLE: map repeating values to rows

```
{
  "movie_id": 200,
  "title": "Alien",
  "sku": "RSD56032",
  "year": 1979,
  "runtime": 117,
  "crew": [
    {"job": "producer",
     "names": ["David Giler",
               "Gordon Carroll"]},
    {"job": "director",
     "names": ["Ridley Scott"]},
    {"job": "screenwriter",
     "names": ["Dan O'Bannon",
               "Ronald Shusett"]}
  ]
}
```

```
select jt.*
from movies m, JSON_TABLE(m.data, '$' columns (
  "ID" path '$.movie_id',
  title,
  year NUMBER,
  nested path '$.crew[*]' columns (
    job,
    nested path '$.names[*]' columns("NAME" path '$'))))
jt;
```

ID	TITLE	YEAR	JOB	NAME
200	Alien	1979	producer	David Giler
200	Alien	1979	producer	Gordon Carroll
200	Alien	1979	director	Ridley Scott
200	Alien	1979	screenwriter	Dan O'Bannon
200	Alien	1979	screenwriter	Ronald Shusett

NESTED is a shortcut for left outer with JSON_TABLE

```
{
  "movie_id": 200,
  "title": "Alien",
  "sku": "RSD56032",
  "year": 1979,
  "runtime": 117,
  "crew": [
    {"job": "producer",
     "names": ["David Giler",
               "Gordon Carroll"]},
    {"job": "director",
     "names": ["Ridley Scott"]},
    {"job": "screenwriter",
     "names": ["Dan O'Bannon",
               "Ronald Shusett"]}
  ]
}
```

```
select jt.*
from movies nested data columns (
  "ID" path '$.movie_id',
  title,
  year NUMBER,
  nested path '$.crew[*]' columns (
    job,
    nested path '$.names[*]' columns("NAME" path '$'))))
jt;
```

ID	TITLE	YEAR	JOB	NAME
200	Alien	1979	producer	David Giler
200	Alien	1979	producer	Gordon Carroll
200	Alien	1979	director	Ridley Scott
200	Alien	1979	screenwriter	Dan O'Bannon
200	Alien	1979	screenwriter	Ronald Shusett

JSON Generation: "from Tables to JSON Data"

Table "customers"

ID	FIRST	LAST	COUNTRY
11	Eliana	Carillo	US
12	Keiran	Stanton	EN
13	Shanice	Collins	US

Table "watches"

cust_id	movie_id	type
11	1	Subscription
11	2	Subscription
12	1	Purchase
13	2	Trial

Table "movie"

ID	TITLE	YEAR	DATA
1	Avatar	2009	{..}
2	Ghostbusters II	1989	{..}

```
select JSON {'name' : c.first || ' ' || c.last,  
            'movies': [select m. title  
                        from movies m, watched w  
                        where w.movie_id = m.id  
                        and w.cust_id = c.id]  
            }  
from customers c;
```

```
{  
  "name": "Eliana Carrilo",  
  "movies": ["Avatar", "Ghostbusters II"]  
}  
  
{  
  "name": "Keiran Stanton", "movies": ["Avatar"]  
}  
  
{  
  "name": "Shanice Collins", "  
  movies": ["Ghostbusters II"]  
}
```

Automatically Map JSON to Rows

```
SELECT JSON_DATAGUIDE(movie_details,  
                      dbms_json.format_hierarchical,  
                      dbms_json.pretty)  
  
FROM   movies;
```

Automatically analyze and present JSON data as "table"

- Schema structure translation to rows and columns
- Instantaneous relational access without any human intervention



Partial updates to JSON documents

```
UPDATE movies m
SET  m.data = json_transform(m.data,
    set    '$.budgetUnit'= 'Million USD',
    set    '$.budget'= (m.data.budget / 1000000),
    remove '$.wiki_article',
    append '$?(@.views > 5000).awards' = 'popularMovie');
```

- Optimized in-place update of JSON documents
- Only update on disk the parts that have been changed
- For larger documents, can provide significant performance improvement over other document stores

Indexing

Single value index –
document indexed once
\$.firstName

```
{
  "firstName": "John" ,
  "lastName": "Smith",
  "age": 25,
  "address": {
    "street": "21 2nd Street", "city": "New York",
    "state": "NY", "postalCode": "10021",
    "isBusiness": false
  },
  "phoneNumbers": [
    { "type": "home", "number": "212 555-1234" },
    { "type": "mobile", "number": "646 555-4567" }
  ],
  "lastUpdated": "2019-05-13T13:03:35+0000"
}
```

Multivalue index –
One document indexed multiple times
\$.phoneNumbers.number



Multi-Value Index – 23c

```
CREATE TABLE jsontab (a NUMBER, doc JSON);
INSERT INTO jsontab VALUES (1,
  '{"credit_score": [710, 720, 710, 730]}');
INSERT INTO jsontab VALUES (2,
  '{"credit_score": [750, 730, 750, 750]}');

CREATE MULTIVALUE INDEX jsontab_idx
  ON jsontab t
  (t.doc.credit_score.number());

SELECT COUNT(*)
FROM jsontab t
WHERE JSON_EXISTS(t.doc,
  '$.credit_score?(@ == 750)');

COUNT(*)
-----
1
```

Multi-Value Index

- Allows to index JSON Array values
- Improved DML performance
- On-par query performance with current existence join plan with parent collection table
- Relies on B-Trees
- 19c workaround: JSON_Table materialized views



Support for JSON in External Tables

- Access JSON files
 - Filesystem
 - Object Storage
- Supports multiple formats automatically
 - Line-delimited
 - Array-wrapped JSON
 - Nested JSON
- Exposed as JSON type column

The screenshot displays the Oracle Database Actions SQL Worksheet interface. The top navigation bar includes the Oracle logo, 'Database Actions | SQL', a search bar, and a 'DUALITY' dropdown. The main workspace shows a SQL script with the following lines:

```
1 CREATE TABLE jsondoc (  
2   data json  
3 )  
4 ORGANIZATION EXTERNAL (  
5   TYPE ORACLE_BIGDATA  
6   DEFAULT DIRECTORY data_pump_dir  
7   ACCESS PARAMETERS(  
8     com.oracle.bigdata.fileformat=jsondoc  
9   )  
10  LOCATION ('b.json')  
11 );  
12  
13 select json_serialize(data) from jsondoc;
```



Below the script, the 'Query Result' tab is active, showing the execution output. The execution time is 0.162 seconds. The result is displayed in a table with one column, 'JSON_SERIALIZE(DATA)', and one row containing the JSON string: '{"a":1,"b":2}'.

	JSON_SERIALIZE(DATA)
1	{"a":1,"b":2}

JSON Collections



Oracle API for MongoDB



MongoDB Application



Oracle Database

- MongoDB developers keep using same skills, tools, and frameworks
- Simplifies migrations from MongoDB to Oracle
- MongoDB does not have tables – it stores collections of JSON documents

JSON Collections

A document is a JSON value

Structure is flexible

A collection contains documents

Supports insert, get, update, filter

A database contains collections

Access data programmatically –
"No SQL"

MongoDB Example

```
MongoClient mongoClient = MongoClient.create(connString);
MongoDatabase database = mongoClient.getDatabase("admin");

MongoCollection<Document> coll =
    database.createCollection("movies");

Document movie = Document.parse(json);
coll.insertOne(movie);

Bson filter = eq("title", "Iron Man");
MongoCursor<Document> cursor = coll.find(filter).cursor();
Document doc = cursor.next();
```

JSON Collections

Database => Schema

Collections created in database "admin" will be in the "ADMIN" schema

MongoDB Example

```
MongoClient mongoClient = MongoClient.create(connString);
MongoDatabase database = mongoClient.getDatabase("admin");

MongoCollection<Document> coll =
    database.createCollection("movies");

Document movie = Document.parse(json);
coll.insertOne(movie);

Bson filter = eq("title", "Iron Man");
MongoCursor<Document> cursor = coll.find(filter).cursor();
Document doc = cursor.next();
```


JSON Collections

Collection => Table

Collections are an abstraction or view of a table with a single JSON column.

```
create table movies
(
  ID VARCHAR2,
  DATA JSON
)
```

MongoDB Example

```
MongoClient mongoClient = MongoClient.create(connString);
MongoDatabase database = mongoClient.getDatabase("admin");

MongoCollection<Document> coll =
    database.createCollection("movies");

Document movie = Document.parse(json);
coll.insertOne(movie);

Bson filter = eq("title", "Iron Man");
MongoCursor<Document> cursor = coll.find(filter).cursor();
Document doc = cursor.next();
```

JSON Collections

Document => Row

Inserting a document into a collection inserts a row into the backing table.

```
insert into  
  movies (data)  
  values (:1);
```

MongoDB Example

```
MongoClient mongoClient = MongoClient.create(connString);  
MongoDatabase database = mongoClient.getDatabase("admin");  
  
MongoCollection<Document> coll =  
  database.createCollection("movies");  
  
Document movie = Document.parse(json);  
coll.insertOne(movie);  
  
Bson filter = eq("title", "Iron Man");  
MongoCursor<Document> cursor = coll.find(filter).cursor();  
Document doc = cursor.next();
```

JSON Collections

Filter => Query

Filter expressions are executed as SQL over the backing table. Fully utilizes core Oracle Database features such as indexing, cost-based optimization, etc.

```
select data
from movies e
where
  e.data.title = 'Iron Man'
```

MongoDB Example

```
MongoClient mongoClient = MongoClient.create(connString);
MongoDatabase database = mongoClient.getDatabase("admin");

MongoCollection<Document> coll =
  database.createCollection("movies");

Document movie = Document.parse(json);
coll.insertOne(movie);

Bson filter = eq("title", "Iron Man");
MongoCursor<Document> cursor = coll.find(filter).cursor();
Document doc = cursor.next();
```


SQL – but only when you need it...

JSON Collections

```
movies.insertOne({  
  "_id" : 123,  
  "title" : "Iron Man"  
});
```

*Simple, flexible persistence
for applications, microservices*

movies



SQL/JSON

```
select t.data.title.string()  
from movies t  
where t.data._id = 123;
```

*Powerful analytics and reporting
directly over collections*

New feature! MongoDB API support for Aggregation Pipelines

The screenshot shows the 'Explain Plan' interface in a MongoDB client. The 'stages' array contains a single stage with a '\$sql' object. The '\$sql' object contains a 'queryPlanner' object and a 'winningPlan' object. The 'winningPlan' object has a 'generatedSql' field that displays the translated Oracle SQL query. The 'executionPlan' field shows the plan hash value. The 'rejectPlans' array is empty.

```

serverInfo: Object
stages: Array
  0: Object
    $sql: Object
      queryPlanner: Object
        plannerVersion: 1
        namespace: "scott.movies"
        indexFilterSet: false
        parsedQuery: Object
        winningPlan: Object
          generatedSql: "with
            'Q1' ('DATA') as (select 'DATA' from 'MOVIES'),
            'Q2' ('DATA') as (
              select 'DATA'
              from 'Q1' q
              where JSON_EXISTS('DATA','$?(@.year.numberOnl
              type(strict))
            ),
            'Q3' ('KEY', 'ACC0') as (
              select
                coalesce(
                  json_query('DATA', '$.year' error on error null c
                  json('null')
                ) as 'KEY',
                json_query('DATA', '$.list_price' error on error n
              from 'Q2' q
            ),
            'Q4' ('DATA') as (
  executionPlan: " Plan Hash Value : 1945855917
rejectPlans: Array
ok: 1
  
```

Output after \$group stage (Sample)

Translated to Oracle analytical SQL

- Mongo pipelines operations equivalent to SQL operations
- Translation to SQL streamlines processing and improves performance

Supports commonly stages and expressions (\$group, \$match, \$project, \$sort, etc.)

Available with ORDS 23.2 for Oracle Database 23c Free later this month



Using the Mongo API with Autonomous Database Serverless

① Enable network access

Access type

~~Secure access from everywhere~~

Allow users with database credentials to access the database from the internet.

Secure access from allowed IPs and VCNs only

Restrict access to specified IP addresses and VCNs. ✓

Private endpoint access only

Restrict access to a private endpoint within an OCI VCN.

Configure access control rules ⓘ

IP notation type Values

IP Address 160.24.114.52

Add my IP address Add access control rule

② Enable the feature

Configure Tools

systems, while the workflows define the sequence in which the data flows are executed. [Learn more.](#)

Web Access (ORDS)	Oracle REST Data Services (ORDS) provides HTTPS interfaces for working with the contents of your Oracle Database in one or more REST enabled schemas. Learn more.	<input checked="" type="checkbox"/> Enabled
MongoDB API	Oracle Database API for MongoDB enables MongoDB-compatible clients and drivers to connect directly to Autonomous Database. Learn more.	<input checked="" type="checkbox"/> Enabled

Apply [Cancel](#)

③ Enable a user

```
create user myuser
  identified by "StrongPass123!";

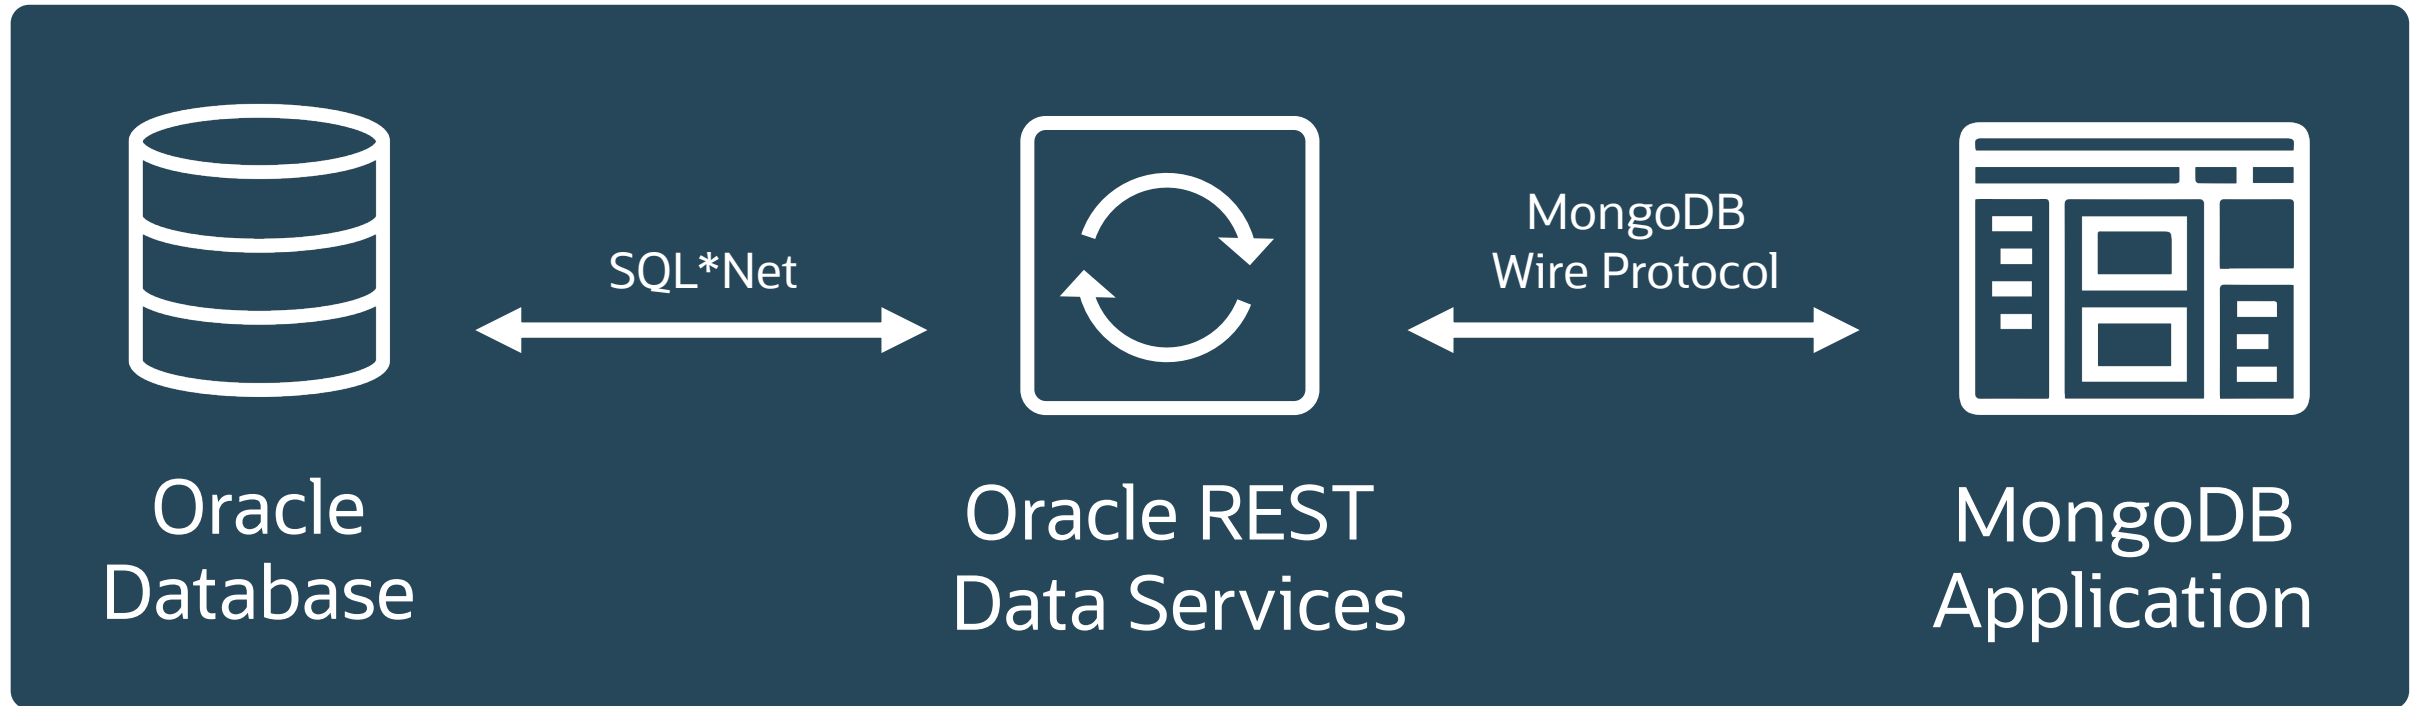
grant dwrole to myuser;
grant unlimited tablespace
  to myuser;

exec ords.enable_schema(
  P_SCHEMA => 'MYUSER'
);
```

On by default for Autonomous JSON



Using the Mongo API with Other Versions of Oracle Database



Autonomous Dedicated, Exadata Database Service, DBaaS, On-premise, etc.



MongoDB API limitations

- MongoAPI uses Oracle database authentication and authorization
 - *use Oracle database username/password only*
- Authenticate as the user corresponding to the database being accessed.
 - *e.g. authenticate as ADMIN, use ADMIN schema/database*
- Indexes must be created from SQL (in release 19c, will be supported in 23c)
 - *db.createIndex() from MongoAPI does not create an index*

```
SQL> create index lastidx on  
      employee (json_value(data, '$.job' error on error));
```


- Aggregation pipeline is only supported on 23c



Demonstration







Citi Bike



<https://www.citibikenyc.com/how-it-works>

Experience NYC in a whole new way

Citi Bike is the nation's largest bike share program, with 14,500 bikes and nearly 900 stations across Manhattan, Brooklyn, Queens and Jersey City. It was designed for quick trips with convenience in mind, and it's a fun and affordable way to get around town.

- 1**  **Join**
- 2**  **Unlock**
- 3**  **Ride**
- 4**  **Return**

SOURCE:

<https://www.citibikenyc.com/>

Bike Station Data



station_information.dat

```
{
  "capacity": 55,
  "eightd_has_key_dispenser": false,
  "eightd_station_services": [],
  "electric_bike_surcharge_waiver": false,
  "external_id": "66db237e-0aca-11e7-82f6-3863bb44ef7c",
  "has_kiosk": true,
  "lat": 40.76727216,
  "lon": -73.99392888,
  "name": "W 52 St & 11 Ave",
  "region_id": "71",
  "rental_methods": [
    "CREDITCARD",
    "KEY"
  ],
  "rental_url": "http://app.citibikenyc.com/S61r/I....",
  "short_name": "6926.01",
  "station_id": "72",
  "station_type": "classic"
}
```

station_status.dat

```
{
  "eightd_has_available_keys": false,
  "is_installed": 1,
  "is_renting": 1,
  "is_returning": 1,
  "last_reported": 1563207770,
  "last_updated": 1563207810,
  "num_bikes_available": 1,
  "num_bikes_disabled": 0,
  "num_docks_available": 52,
  "num_docks_disabled": 2,
  "num_ebikes_available": 0,
  "station_id": "72"
}
```

N

1

Oracle's Strategy on JSON



Leader: first enterprise relational database with JSON support: 12.1.0.2 (2014)

- large number of customers, NoSQL migrations, 26% of Autonomous DB use JSON

Open: JSON enhancement are contributed to the open standards

- SQL standard: SAP Hana, IBM DB2, Postgres, MySQL,... follow Oracle
- JSON Schema: Oracle wrote the database vocabulary

Innovative: many new industry leading JSON features

- JSON_Dataguide to derive JSON schema from data, detect schema changes, auto create relational views
- JSON Schema to validate JSON but also describe database objects, constraints
- JSON Relational Duality, GraphQL, ...



Thank you

